

Introduction to Data Analysis with Pandas

Contents

- 1.1. Data description
- 1.2. Loading data
- 1.3. Visualize single variable data
- 1.4. Scatter plots for multiple variables
- 1.5. Sort the data
- 1.6. Describe categorical data
- 1.7. Filter and split data
- 1.8. Creating new data frames and adding new columns to data frames
- 1.9. Box plot of numerical data sorted by category
- 1.10. Multi-plot figures
- 1.11. Linear Regression
- 1.12. The Hertzsprung-Russell Diagram

[Pandas](#) is a library for data analysis, manipulation, and visualization. The basic object of the defined by this module is the [DataFrame](#). This is a dataset used in this notebook can be obtained from Kaggle on the [classification of stars](#). We load the data from a CSV file into a Pandas DataFrame and demonstrate some basic functionality of the module.

You can think of **data frames** as tables basically, where each row is an data entry and each of the columns is a property of that entry. In our case, each entry is gonna be a star and the columns some of its properties. We could also use each row to store the properties of some system at different time-points, eg. the concentration of different proteins over time in a cell, in this case each row would be a time-point and each column would be the different proteins.

As opposed to numpy arrays, Pandas data frames allow to work with the data by using labels —eg. 'temperature'— rather than having to remember the index numbers

[Skip to main content](#)

Another great aspect of Pandas data frames is that we can mix types of data, eg. numerical variables like the mass of object —eg. 21.2 mg— and categorical data like a cell type —eg. 'cortical neuron'—.

In the field of data science the columns of dataset are often referred as 'feature vectors'. If you encounter that term, simply replace it in your mind by 'column'.

1.1. Data description

Each row represent a star.

Feature vectors:

- Temperature – The surface temperature of the star
- Luminosity – Relative luminosity: how bright it is
- Size – Relative radius: how big it is
- AM – [Absolute magnitude](#): another measure of the star luminosity
- Color – General Color of Spectrum
- Type – Red Dwarf, Brown Dwarf, White Dwarf, Main Sequence , Super Giants, Hyper Giants
- Spectral_Class – O,B,A,F,G,K,M / SMASS - [Stellar classification](#)

1.2. Loading data

```
import numpy as np
import matplotlib.pyplot as plt
import matplotlib
import pandas as pd
pd.options.mode.chained_assignment = None # default='warn'
```

/var/folders/5r/37slq4k14kb2dywqvcdlbwdc0000gn/T/ipykernel_71216/2359310441.py:4: D
Pyarrow will become a required dependency of pandas in the next major release of pa
(to allow more performant data types, such as the Arrow string type, and better int
but was not found to be installed on your system.

If this would cause problems for you,
please provide us feedback at <https://github.com/pandas-dev/pandas/issues/54466>

```
import pandas as pd
```

The DataFrame can be created from a csv file using the `read_csv` method. If you are working on Colab, you will need to upload the data.

Notice that in this case we are loading data from a .csv file, but with Pandas we can load pretty much any kind of data format, including matlab data files.

```
df = pd.read_csv('Stars.csv')
```

The head method displays the first few rows of data together with the column headers

```
df.head()
```

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
0	3068	0.002400	0.1700	16.12	Red	M	Red Dwarf
1	3042	0.000500	0.1542	16.60	Red	M	Red Dwarf
2	2600	0.000300	0.1020	18.70	Red	M	Red Dwarf
3	2800	0.000200	0.1600	16.65	Red	M	Red Dwarf
4	1939	0.000138	0.1030	20.06	Red	M	Red Dwarf

Specific columns of the DataFrame can be accessed by specifying the column name between square brackets:

```
stars_colors = df['Color'] # notice the columns names are case sensitive, ie. 'color'
print(stars_colors)
```

```
0      Red
1      Red
2      Red
3      Red
4      Red
...
235    Blue
236    Blue
237    White
238    White
239    Blue
Name: Color, Length: 240, dtype: object
```

The individual entries of the DataFrame (ie. rows) can be accessed using the `iloc`

[Skip to main content](#)

```
print(df.iloc[[0]]) # where 0 is the index of the first entry
```

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
0	3068	0.0024	0.17	16.12	Red	M	Red Dwarf

The describe method will give basic summary statistics for the numerical variables of each column

```
summary = df.describe()
print(summary)
```

	Temperature	Luminosity	Size	A_M
count	240.000000	240.000000	240.000000	240.000000
mean	10497.462500	107188.361635	237.157781	4.382396
std	9552.425037	179432.244940	517.155763	10.532512
min	1939.000000	0.000080	0.008400	-11.920000
25%	3344.250000	0.000865	0.102750	-6.232500
50%	5776.000000	0.070500	0.762500	8.313000
75%	15055.500000	198050.000000	42.750000	13.697500
max	40000.000000	849420.000000	1948.500000	20.060000

We can also call methods of the individual columns to get summary information. The column objects (such as df['Temperature']) are called Series

```
print("Mean Temperature is:",df['Temperature'].mean())
print("Max Temperature is:",df['Temperature'].max())
```

```
Mean Temperature is: 10497.4625
Max Temperature is: 40000
```

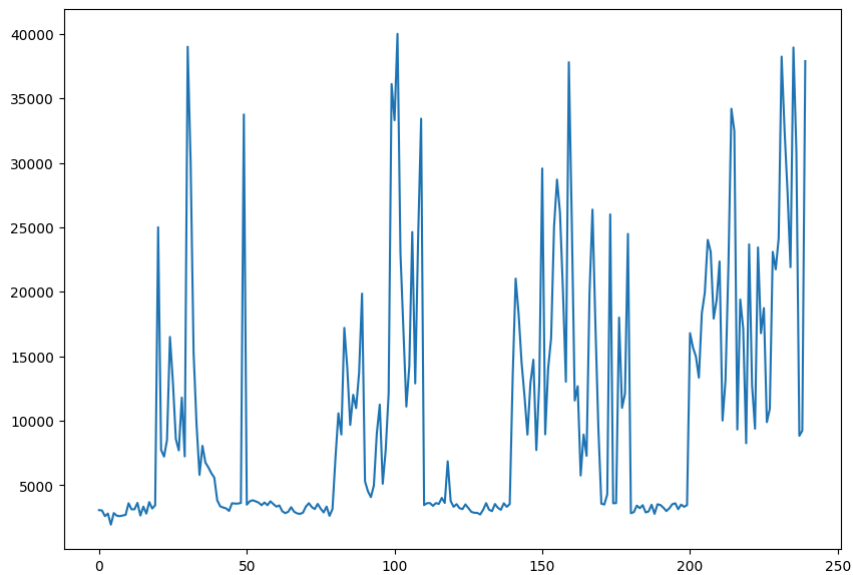
1.3. Visualize single variable data

The Series objects (columns) have plot methods as well as the numerical summary methods.

```
df['Temperature'].plot.line(figsize=(10,7))
```

✓ Aves · \

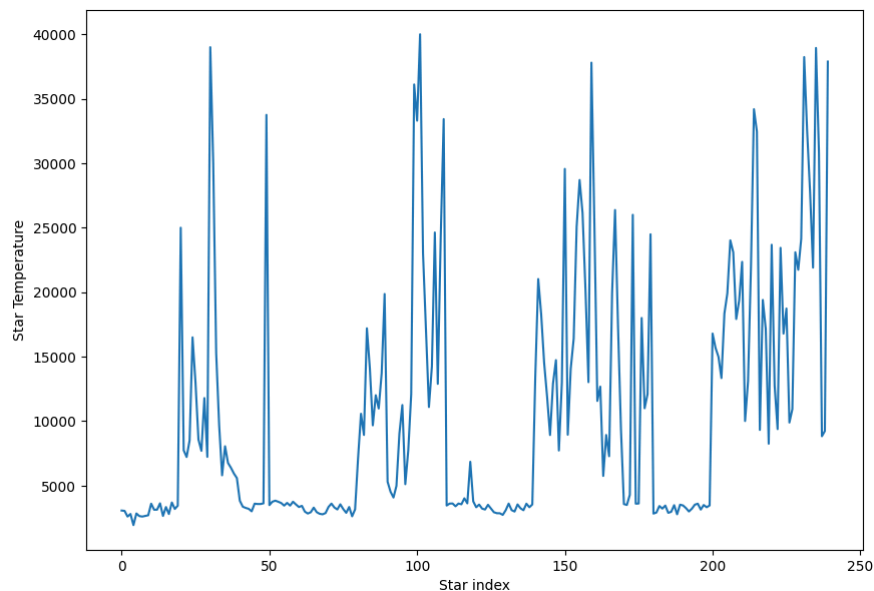
[Skip to main content](#)



Pandas is interoperable with matplotlib and numpy, so for instance if we want to add labels to the figure above we simply add the following lines from matplotlib:

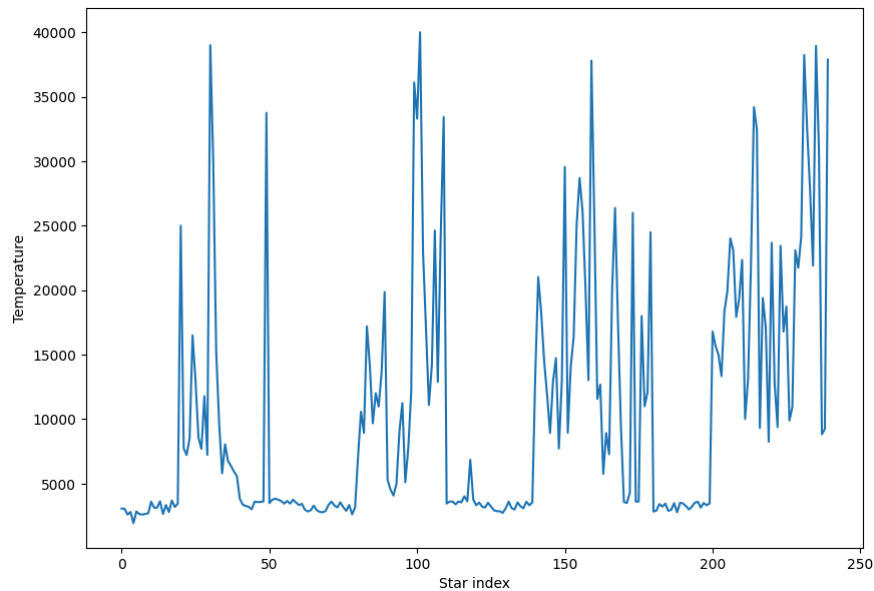
```
df['Temperature'].plot.line(figsize=(10,7));  
plt.xlabel('Star index')  
plt.ylabel('Star Temperature')
```

```
Text(0, 0.5, 'Star Temperature')
```



The above is equivalent to:

```
df['Temperature'].plot.line(xlabel = 'Star index', ylabel='Temperature', figsize=(10,7))
```



1.3.1. Exercise

- Check Pandas [series.plot documentation](#) and plot the temperature of the different stars as an histogram.
- By observing at the histogram, what's the most common temperature for stars?

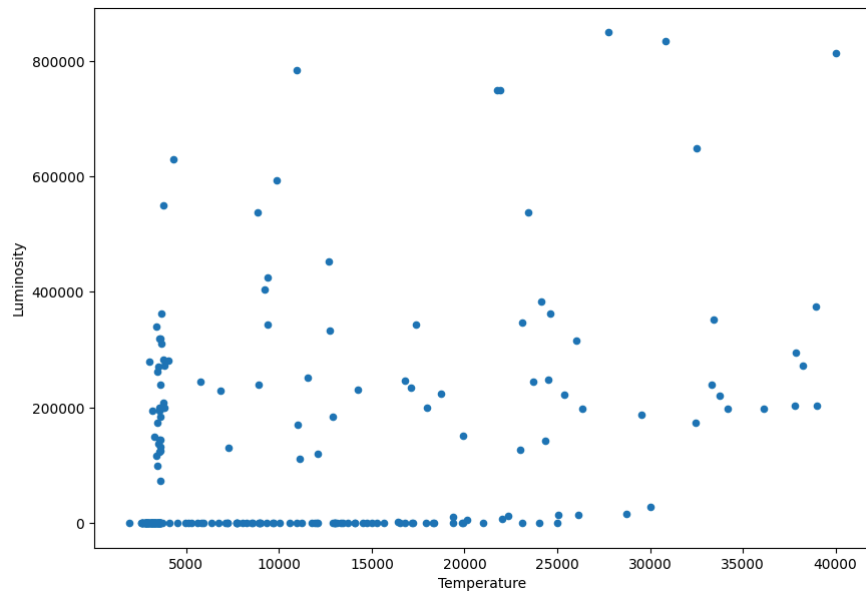
```
## Your code here
```

1.4. Scatter plots for multiple variables

A typical problem in any field is to understand how some properties relate others, eg. are two properties independent or correlated to each other? We can quickly explore the correlations in some data frame by using scatter plots and plotting some properties against others:

```
df.plot.scatter('Temperature', 'Luminosity', figsize=(10,7))
```

```
<Axes: xlabel='Temperature', ylabel='Luminosity'>
```



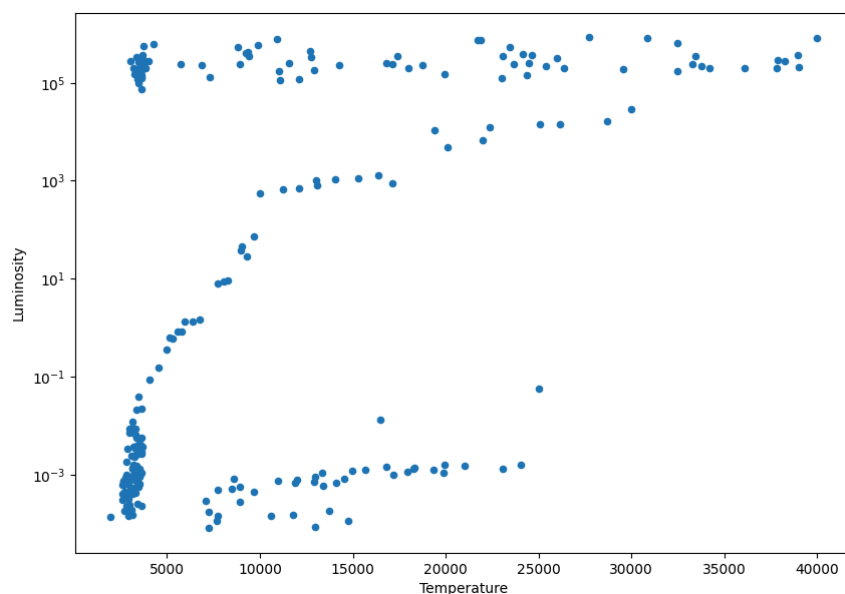
We notice that the values of the Luminosity go from very small to very big values:

```
print(df['Luminosity'].min())  
print(df['Luminosity'].max())
```

```
8e-05  
849420.0
```

In this situations where we are plotting over a very long range of values, it's useful to change the scale to a logarithmic one:

```
ax = df.plot.scatter('Temperature', 'Luminosity', figsize=(10,7))  
plt.yscale('log')
```



[Skip to main content](#)

1.4.1. Exercise

- Make scatter plots of the different star features.
- Two of the feature columns in the data are monotonically correlated, find them. #
Hint: you may need to use log scale to better see a linear correlation.

```
## Your code here
```

1.5. Sort the data

We can sort the data using the `sort_values` method:

```
sorted_data = df.sort_values('Temperature', ascending=True)
sorted_data.head()
```

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
4	1939	0.000138	0.103	20.06	Red	M	Red Dwarf
2	2600	0.000300	0.102	18.70	Red	M	Red Dwarf
7	2600	0.000400	0.096	17.40	Red	M	Red Dwarf
78	2621	0.000600	0.098	12.81	Red	M	Brown Dwarf
6	2637	0.000730	0.127	17.22	Red	M	Red Dwarf

1.6. Describe categorical data

We can describe the categorical variable 'Color'. In this case we get different results than when we used describe on a numerical value.

```
print(df['Color'].describe())
```

```
count      240
unique       17
top         Red
freq        112
Name: Color, dtype: object
```

[Skip to main content](#)


```
print(df['Color'].unique())
```

```
['Red' 'Blue White' 'White' 'Yellowish White' 'Blue white'  
'Pale yellow orange' 'Blue' 'Blue-white' 'Whitish' 'yellow-white'  
'Orange' 'White-Yellow' 'white' 'yellowish' 'Yellowish' 'Orange-Red'  
'Blue-White']
```

1.6.1. Exercise

- Create a histogram to visualize how many stars of each color there are.

```
## Your code here
```

1.7. Filter and split data

Sometimes we want to select sections of the data based on their values, we can easily do so with pandas. Let's find the set of stars whose temperature is higher than 10000 K. We first create a *boolean array* for the condition, that is, a vector which associate a true or false value to each star with regard to the filtering condition, in our case case, it will give a true value if the start temperature is higher than 10000 K and false otherwise:

```
hot_stars_boolean_vector = df['Temperature'] > 10000  
print(hot_stars_boolean_vector)
```

```
0      False  
1      False  
2      False  
3      False  
4      False  
...  
235     True  
236     True  
237     False  
238     False  
239     True  
Name: Temperature, Length: 240, dtype: bool
```

In python a true value is represented with the number 1 and a false value with the number zero, that means that if we want to know hot many stars are hotter than 10000 K

[Skip to main content](#)

```
number_of_hot_stars = np.sum(hot_stars_boolean_vector)
print(f'There are {number_of_hot_stars} hot stars in the dataset')
```

There are 90 hot stars in the dataset

It works the same for categorical data. Let's find out the number of super giants stars:

```
super_giants_boolean_vector = df['Type'] == 'Super Giants'
nb_super_giants = np.sum(super_giants_boolean_vector)
print(f'There are {nb_super_giants} super giants stars in the dataset')
```

There are 40 super giants stars in the dataset

If we are only interested in exploring the properties of super giants stars (because our dataset is too big or because white dwarfs are lame), we can get select only the data of the super giants stars using the boolean vector we just created:

```
df_with_only_super_giants = df[super_giants_boolean_vector]
print(df_with_only_super_giants)
```

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
40	3826	200000.0	19.0	-6.930	Red	M	Super Giants
41	3365	340000.0	23.0	-6.200	Red	M	Super Giants
42	3270	150000.0	88.0	-6.020	Red	M	Super Giants
43	3200	195000.0	17.0	-7.220	Red	M	Super Giants
44	3008	280000.0	25.0	-6.000	Red	M	Super Giants
45	3600	320000.0	29.0	-6.600	Red	M	Super Giants
46	3575	123000.0	45.0	-6.780	Red	M	Super Giants
47	3574	200000.0	89.0	-5.240	Red	M	Super Giants
48	3625	184000.0	84.0	-6.740	Red	M	Super Giants
49	33750	220000.0	26.0	-6.100	Blue	B	Super Giants
100	33300	240000.0	12.0	-6.500	Blue	B	Super Giants
101	40000	813000.0	14.0	-6.230	Blue	O	Super Giants
102	23000	127000.0	36.0	-5.760	Blue	O	Super Giants
103	17120	235000.0	83.0	-6.890	Blue	O	Super Giants
104	11096	112000.0	12.0	-5.910	Blue	O	Super Giants
105	14245	231000.0	42.0	-6.120	Blue	O	Super Giants
106	24630	363000.0	63.0	-5.830	Blue	O	Super Giants
107	12893	184000.0	36.0	-6.340	Blue	O	Super Giants
108	24345	142000.0	57.0	-6.240	Blue	O	Super Giants
109	33421	352000.0	67.0	-5.790	Blue	O	Super Giants
160	25390	223000.0	57.0	-5.920	Blue	O	Super Giants
161	11567	251000.0	36.0	-6.245	Blue	O	Super Giants
162	12675	452000.0	83.0	-5.620	Blue	O	Super Giants
163	5752	245000.0	97.0	-6.630	Blue	O	Super Giants

[Skip to main content](#)

166	19923	152000.0	73.0	-5.690	Blue	0	Super Giants
167	26373	198000.0	39.0	-5.830	Blue	0	Super Giants
168	17383	342900.0	30.0	-6.090	Blue	0	Super Giants
169	9373	424520.0	24.0	-5.990	Blue	0	Super Giants
220	23678	244290.0	35.0	-6.270	Blue	0	Super Giants
221	12749	332520.0	76.0	-7.020	Blue	0	Super Giants
222	9383	342940.0	98.0	-6.980	Blue	0	Super Giants
223	23440	537430.0	81.0	-5.975	Blue	0	Super Giants
224	16787	246730.0	62.0	-6.350	Blue	0	Super Giants
225	18734	224780.0	46.0	-7.450	Blue	0	Super Giants
226	9892	593900.0	80.0	-7.262	Blue	0	Super Giants
227	10930	783930.0	25.0	-6.224	Blue	0	Super Giants
228	23095	347820.0	86.0	-5.905	Blue	0	Super Giants
229	21738	748890.0	92.0	-7.346	Blue	0	Super Giants

Wait, wait. What if we want we want to filter for two conditions, say, we want to keep only the very hoy super giant stars? Low and behold, we simply need to apply both conditions:

```
hot_super_giants = df[super_giants_boolean_vector & hot_stars_boolean_vector]
print(f"There are {hot_super_giants.shape[0]} super hot giants")
print(hot_super_giants)
```

There are 25 super hot giants

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
49	33750	220000.0	26.0	-6.100	Blue	B	Super Giants
100	33300	240000.0	12.0	-6.500	Blue	B	Super Giants
101	40000	813000.0	14.0	-6.230	Blue	0	Super Giants
102	23000	127000.0	36.0	-5.760	Blue	0	Super Giants
103	17120	235000.0	83.0	-6.890	Blue	0	Super Giants
104	11096	112000.0	12.0	-5.910	Blue	0	Super Giants
105	14245	231000.0	42.0	-6.120	Blue	0	Super Giants
106	24630	363000.0	63.0	-5.830	Blue	0	Super Giants
107	12893	184000.0	36.0	-6.340	Blue	0	Super Giants
108	24345	142000.0	57.0	-6.240	Blue	0	Super Giants
109	33421	352000.0	67.0	-5.790	Blue	0	Super Giants
160	25390	223000.0	57.0	-5.920	Blue	0	Super Giants
161	11567	251000.0	36.0	-6.245	Blue	0	Super Giants
162	12675	452000.0	83.0	-5.620	Blue	0	Super Giants
166	19923	152000.0	73.0	-5.690	Blue	0	Super Giants
167	26373	198000.0	39.0	-5.830	Blue	0	Super Giants
168	17383	342900.0	30.0	-6.090	Blue	0	Super Giants
220	23678	244290.0	35.0	-6.270	Blue	0	Super Giants
221	12749	332520.0	76.0	-7.020	Blue	0	Super Giants
223	23440	537430.0	81.0	-5.975	Blue	0	Super Giants
224	16787	246730.0	62.0	-6.350	Blue	0	Super Giants
225	18734	224780.0	46.0	-7.450	Blue	0	Super Giants
227	10930	783930.0	25.0	-6.224	Blue	0	Super Giants
228	23095	347820.0	86.0	-5.905	Blue	0	Super Giants
229	21738	748890.0	92.0	-7.346	Blue	0	Super Giants

```
hot_super_giants = df[(df['Type'] == 'Super Giants') & (df['Temperature'] > 10000)]
print(f"\nThere are {hot_super_giants.shape[0]} super hot giants\n")
print(hot_super_giants)
```

There are 25 super hot giants

	Temperature	Luminosity	Size	A_M	Color	Spectral_Class	Type
49	33750	220000.0	26.0	-6.100	Blue	B	Super Giants
100	33300	240000.0	12.0	-6.500	Blue	B	Super Giants
101	40000	813000.0	14.0	-6.230	Blue	O	Super Giants
102	23000	127000.0	36.0	-5.760	Blue	O	Super Giants
103	17120	235000.0	83.0	-6.890	Blue	O	Super Giants
104	11096	112000.0	12.0	-5.910	Blue	O	Super Giants
105	14245	231000.0	42.0	-6.120	Blue	O	Super Giants
106	24630	363000.0	63.0	-5.830	Blue	O	Super Giants
107	12893	184000.0	36.0	-6.340	Blue	O	Super Giants
108	24345	142000.0	57.0	-6.240	Blue	O	Super Giants
109	33421	352000.0	67.0	-5.790	Blue	O	Super Giants
160	25390	223000.0	57.0	-5.920	Blue	O	Super Giants
161	11567	251000.0	36.0	-6.245	Blue	O	Super Giants
162	12675	452000.0	83.0	-5.620	Blue	O	Super Giants
166	19923	152000.0	73.0	-5.690	Blue	O	Super Giants
167	26373	198000.0	39.0	-5.830	Blue	O	Super Giants
168	17383	342900.0	30.0	-6.090	Blue	O	Super Giants
220	23678	244290.0	35.0	-6.270	Blue	O	Super Giants
221	12749	332520.0	76.0	-7.020	Blue	O	Super Giants
223	23440	537430.0	81.0	-5.975	Blue	O	Super Giants
224	16787	246730.0	62.0	-6.350	Blue	O	Super Giants
225	18734	224780.0	46.0	-7.450	Blue	O	Super Giants
227	10930	783930.0	25.0	-6.224	Blue	O	Super Giants
228	23095	347820.0	86.0	-5.905	Blue	O	Super Giants
229	21738	748890.0	92.0	-7.346	Blue	O	Super Giants

1.7.1. Exercise

- Find how many 'White Dwarf' have a surface temperature between 5000 K and 10000 K
- Find the mean surface temperature of the White Dwarfs
- How many times bigger are Super Giants stars compared to White Dwarfs?
- What's the variance in the size of Super Giant stars?

```
## Your code here
```

1.8. Creating new data frames and adding new columns to data frames

We can create a new data frame from another one with only some of the original data frame columns. Let's create a new data frame with only the temperature and type columns:

```
new_df = df[['Temperature', 'Type']]
print(new_df.head()) # It's always good practice to print the head of the data frame
```

	Temperature	Type
0	3068	Red Dwarf
1	3042	Red Dwarf
2	2600	Red Dwarf
3	2800	Red Dwarf
4	1939	Red Dwarf

We may also want to add new columns to an existing data frame, for instance, if we incorporate new data from a different file or we calculate new quantities based on the previous data. Here we are adding a new column whose values are the inverse of the luminosity:

```
df['Inverse Luminosity'] = 1 / df['Luminosity']
```

1.8.1. Exercise

- Add a new feature vector to the new data frame with the volume of each star. # Hint: Notice the column 'Size' is the radius R of each star and that the volume of a sphere is $\frac{4}{3}\pi R^3$
- (Bonus Exercise) Add a new feature vector to the new data frame with the mass of each star. # Hint: The mass m of an object is equal to the product of the volume V by its density ρ , that is, $m = \rho V$. Notice that different types of stars have different densities so you'll have to use the filtering as we did above: $\rho_{Dwarfs} = 10^5 g/cc$, $\rho_{Giants} = 10^{-8} g/cc$, $\rho_{Main\ sequence} = 1 g/cc$. You are welcome to ignore the units, the goal is that you practice how to apply operations to a subset of data frame.

```
## Your code here
```

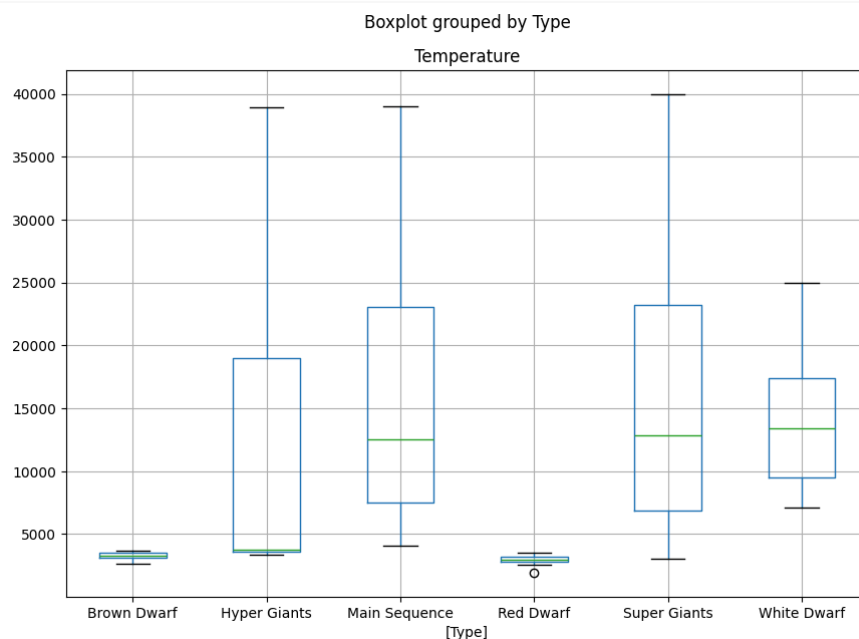
[Skip to main content](#)

1.9. Box plot of numerical data sorted by category

Let's visualise what is the range of temperature of the different stars based on their temperature. To do so, we first select the features we want to visualise and then call a box plot:

```
# The boxplot argument 'by' will split the plot over the variable given.
df[['Temperature', 'Type']].boxplot(by='Type', figsize=(10,7))
```

```
<Axes: title={'center': 'Temperature'}, xlabel='[Type]'
```



1.9.1. Exercise

- Make a similar figure as the above but displaying the range of volumes of the different star types

```
## Your code here
```

1.10. Multi-plot figures

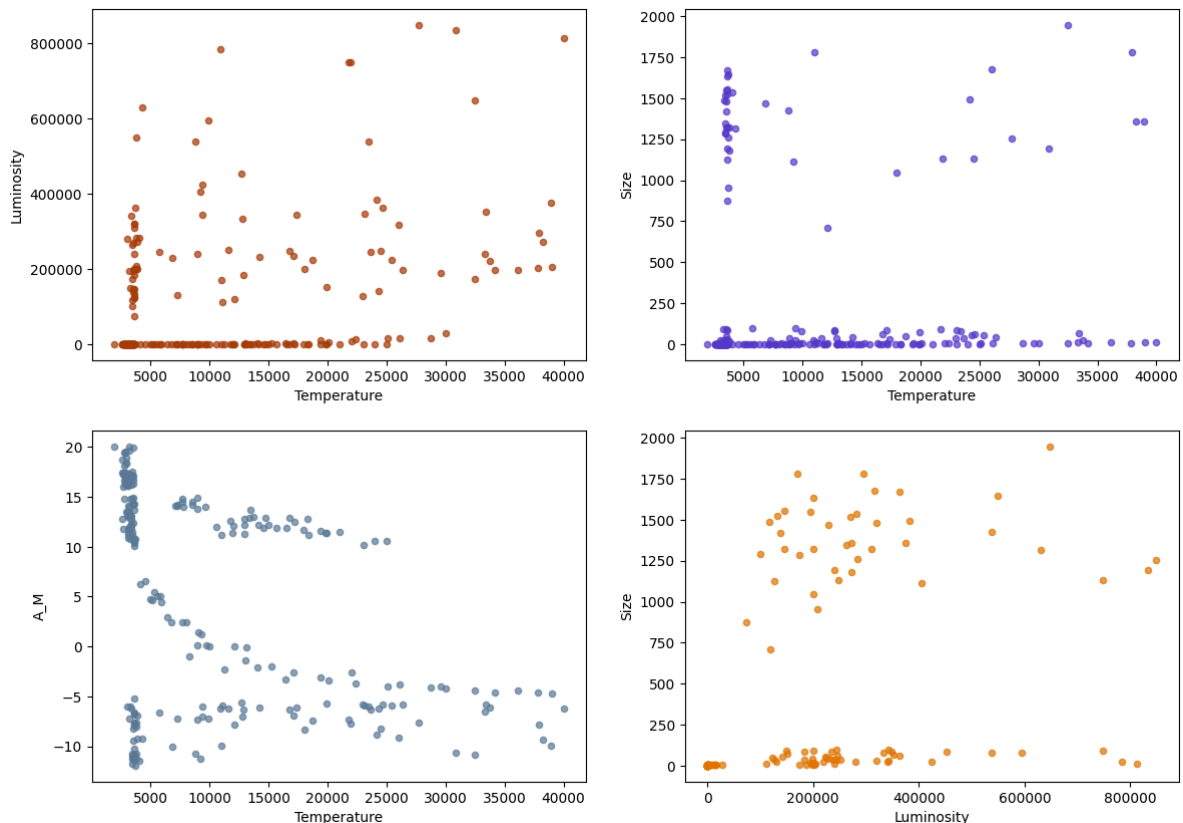
Now that we know how to filter data, let's make some figures. We construct a figure with 4 subplots:

[Skip to main content](#)

```
fig, ax = plt.subplots(2,2, figsize=(16, 12))
fig.set_figwidth(14)
fig.set_figheight(10)

## Plot the data. ax[i,j] references the the Axes in row i column j
df.plot.scatter('Temperature', 'Luminosity', color='xkcd:rust', alpha=0.7, ax=ax[0,0])
df.plot.scatter('Temperature', 'Size', color='xkcd:blue', alpha=0.7, ax=ax[0,1])
df.plot.scatter('Temperature', 'A_M', color='xkcd:slate blue', alpha=0.7, ax=ax[1,0])
df.plot.scatter('Luminosity', 'Size', color='xkcd:pumpkin', alpha=0.7, ax=ax[1,1])
```

<Axes: xlabel='Luminosity', ylabel='Size'>



We can see in the plot of A_M versus Temperature, that there is a cluster of points ($A_M > 9$, Temperature > 5000) where the variables appear to have a strong correlation. We might want to isolate and study that particular subset of the data by extracting it to a different DataFrame.

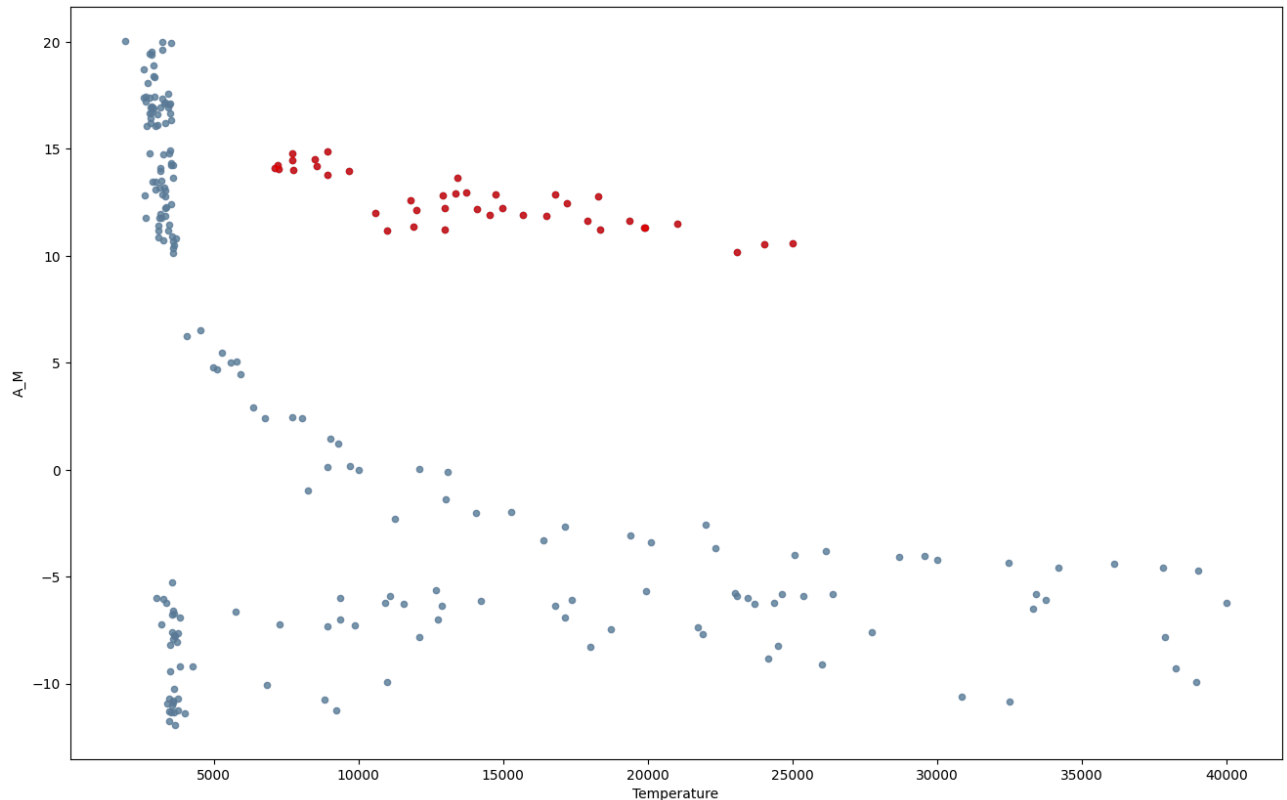
Let's isolate it into the variable `df_TAM` and plot it in a different color:

```
df_AM = df[df['A_M'] > 9]
df_TAM = df_AM[df_AM['Temperature'] > 5000]

## Plot the subset with the original
ax = df.plot.scatter('Temperature', 'A_M', color='xkcd:slate blue', alpha=0.8)
df_TAM.plot.scatter('Temperature', 'A_M', color='xkcd:red', ax=ax, alpha=0.7)
```

[Skip to main content](#)

```
<Axes: xlabel='Temperature', ylabel='A_M'>
```



Let's print the statistics of this subset of data

```
print(df_TAM.describe())
```

	Temperature	Luminosity	Size	A_M	Inverse Luminosity
count	40.000000	40.000000	40.000000	40.000000	40.000000
mean	13931.450000	0.002434	0.010728	12.582500	2836.282072
std	4957.655189	0.008912	0.001725	1.278386	3270.623635
min	7100.000000	0.000080	0.008400	10.180000	17.857143
25%	9488.750000	0.000287	0.009305	11.595000	814.754098
50%	13380.000000	0.000760	0.010200	12.340000	1316.701317
75%	17380.000000	0.001227	0.012025	13.830000	3479.064039
max	25000.000000	0.056000	0.015000	14.870000	12500.000000

1.11. Linear Regression

Let's finish this notebook by doing a linear regression on the data.

A linear regression consists in a linear model that relates one variable to another variable. For instance, the temperature of a star to its luminosity. Linear models have the advantage of being easily interpreted—you can look at the model and figure out what's going on.

[Skip to main content](#)

of making natural phenomena *very* non-linear. On the Machine learning notebooks, we'll learn how to train models that can deal with non-linear dependencies.

Since the data in the high absolute magnitude A_M , high-Temperature subset seem to be strongly correlated, we might fit linear model. To do this we will import the **linregress** function from the **stats** module in SciPy.

```
from scipy.stats import linregress

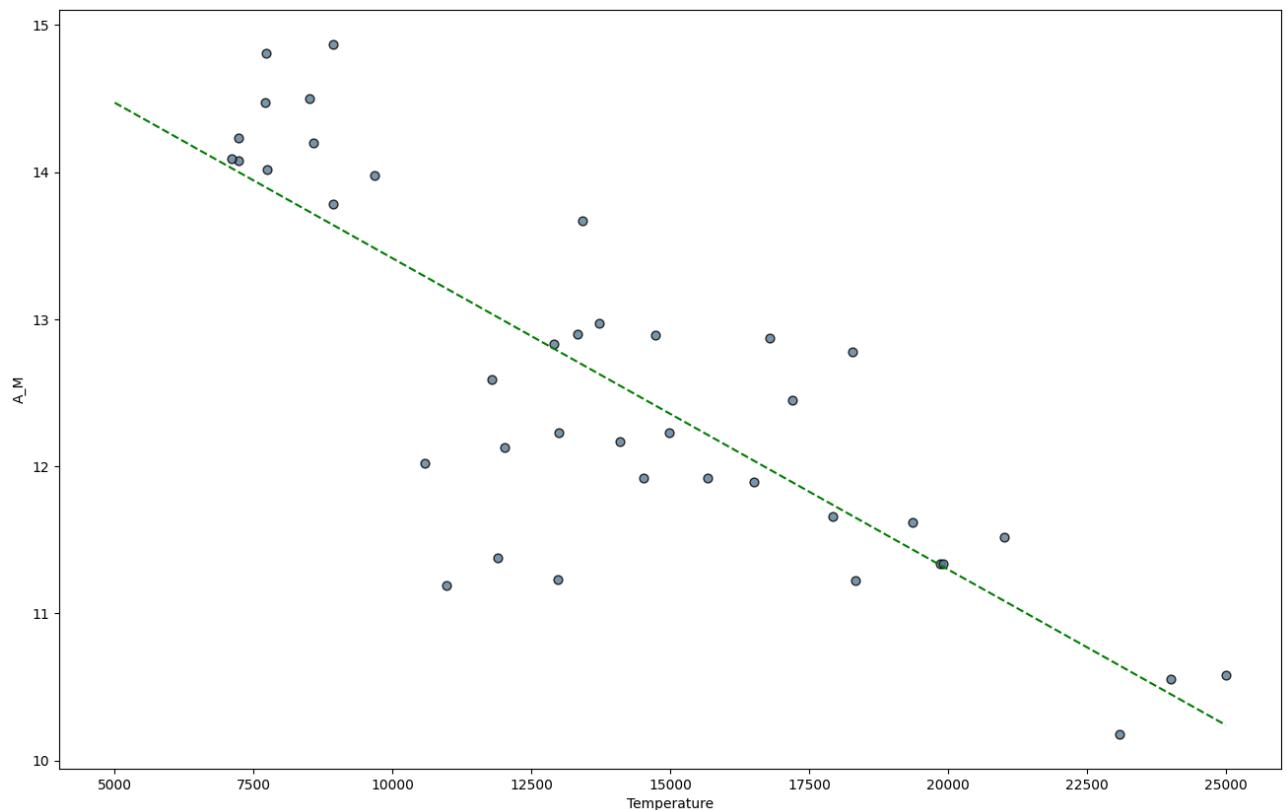
linear_model = linregress(df_TAM['Temperature'], df_TAM['A_M'])
```

Let's plot the regression line together with the data.

```
m = linear_model.slope
b = linear_model.intercept

x = np.linspace(5000, 25000, 5) # Range of temperatures
y = m*x + b
ax = df_TAM.plot.scatter('Temperature', 'A_M', color='xkcd:slate blue', s=40, edgecolor='black')
ax.plot(x, y, color='green', ls='dashed')
```

[<matplotlib.lines.Line2D at 0x17f2afe10>]



The model object that was produced by **linregress** also contains the correlation

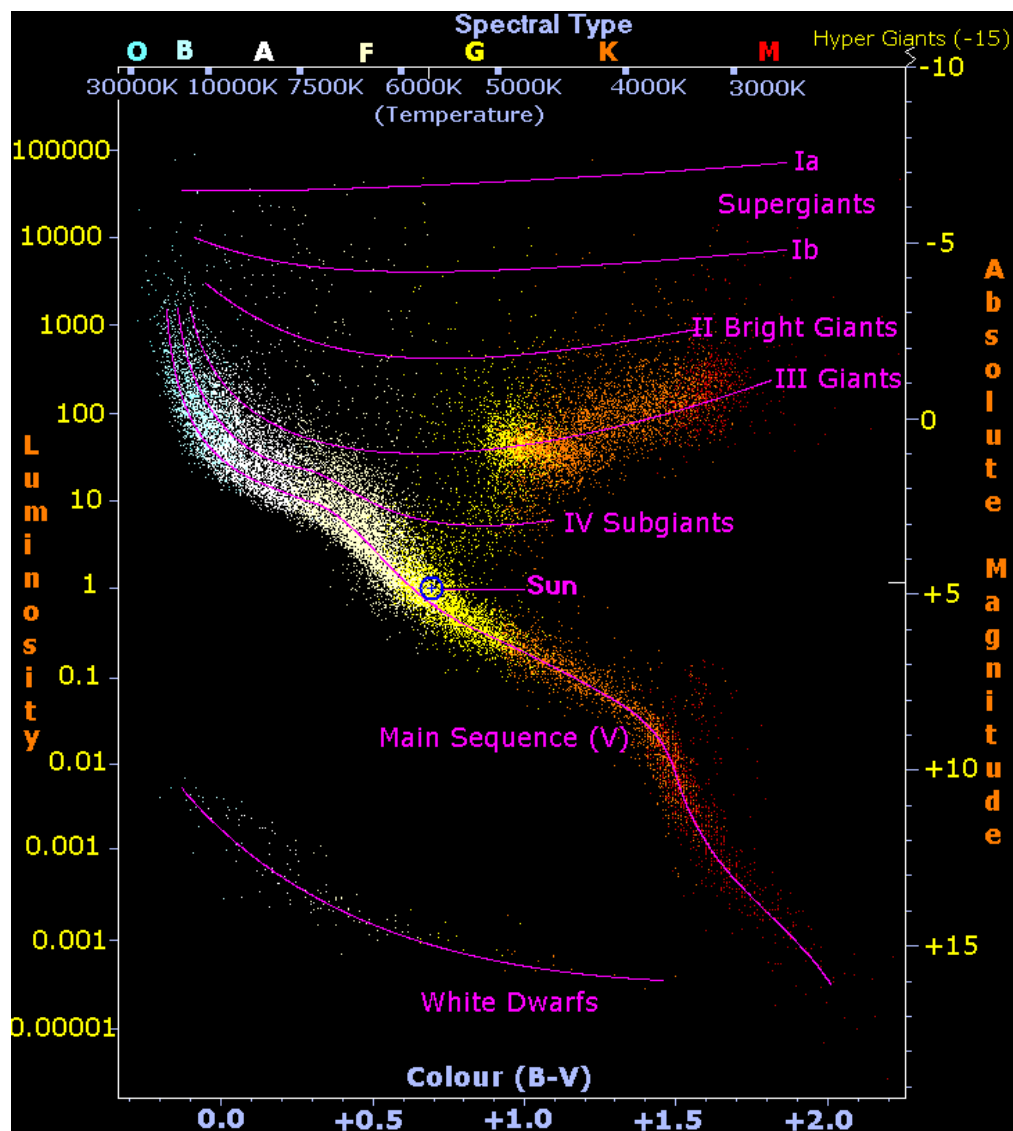
[Skip to main content](#)

```
print("Correlation coefficient:",linear_model.rvalue)
print("pvalue for null hypothesis of slope = 0:",linear_model.pvalue)
print("Standard error of the esimated gradient:",linear_model.stderr)
```

Correlation coefficient: -0.8201933172733418
pvalue for null hypothesis of slope = 0: 9.411965493687398e-11
Standard error of the esimated gradient: 2.3930706947460813e-05

1.12. The Hertzsprung-Russell Diagram

The Hertzsprung-Russell Diagram is a scatter plot of stars showing the relationship between the stars' absolute magnitudes or luminosities versus their temperatures.



Let's see if we can obtain something similar from our data:

1.12.1. Exercise

- Make a scatter plot from our star data. Plot each star type 'Super Giants', 'Main sequence' and 'White Dwarf' in different colours.
- Can you observe similar star clusters? # Hint: You might need to use logarithmic scales for the axis and reverse the direction of the temperature axis.

```
## Your code here
```

Previous

< [4. Introduction to Brian part 3: Simulations](#)

Next

[2. Machine Learning](#) >